

# Unilever POC - Architecture Overview

## Unilever Procurement GPT POC - Architecture Overview

### System Overview

#### 1. Architecture Components

Existing Layer (Unilever)

Data Collection Module

Evaluation Framework (Feature 1)

Monitoring Framework (Feature 2)

#### 2. Storage Layer

Table 1: queries

Table 2: evaluations

Table 3: errors

Table 4: drift\_monitoring

Table 5: baseline

#### 3. Data Flow

# Unilever Procurement GPT POC

## - Architecture Overview

---

### System Overview

---

**Goal:** Evaluate and monitor AI agents (Spend & Demand) with  $\geq 90\%$  accuracy

**Key Components:** 1. Evaluation Framework - Validates agent response correctness 2. Monitoring Framework - Detects drift and classifies errors

---

### 1. Architecture Components

---

#### Existing Layer (Unilever)

- **Spend Agent** - Handles spend queries, generates SQL
- **Demand Agent** - Handles demand queries, generates SQL

#### Data Collection Module

- Captures: Query, Response, Logs, Timestamp
- Format: JSON

#### Evaluation Framework (Feature 1)

##### 6-Step Process:

**Step 1: Pre-Processing** - Clean data, extract SQL, normalize format

**Step 2: Structural Validation** - Check SQL syntax, schema, data types

**Step 3: Semantic Check** - Compare with ground truth, calculate similarity (0.0-1.0)

**Step 4: LLM Judge** - Model: Ollama Llama 3.1 (free, local) - Input: Query + Response + Ground Truth - Output: PASS/FAIL + Confidence + Reasoning

**Step 5: Scoring** - Formula:  $(0.3 \times \text{Structural}) + (0.3 \times \text{Semantic}) + (0.4 \times \text{LLM})$  - Decision: PASS if score  $\geq 0.7$

**Step 6: Store Result** - Save to Evaluation DB with all scores

## Monitoring Framework (Feature 2)

### Layer 1: Query Monitoring

- **Purpose:** Track all incoming queries
- **Process:** Capture query patterns, frequency, complexity
- **Output:** Query statistics for baseline comparison

### Layer 2: Drift Detection

- **Purpose:** Identify when queries deviate from expected patterns
- **Process:**
  1. Convert query to 384-dim vector (embedding)
  2. Compare with baseline (1000 training queries centroid)
  3. Calculate drift score:  $1 - \text{cosine\_similarity}$
  4. Classify: Low (0.1-0.3), Medium (0.3-0.5), High ( $>0.5$ )
- **Trigger:** Alert if High drift detected

### Layer 3: Error Detection

- **Purpose:** Identify failures in agent responses
- **Triggers:**
  - Evaluation FAIL result
  - API exceptions/timeouts
  - SQL execution errors
- **Output:** Error event with full context

### Layer 4: Error Classification

- **Purpose:** Categorize errors for root cause analysis

- **Categories:**

1. SQL Generation - syntax, schema, logic errors
2. Context Retrieval - missing or wrong context
3. Data Errors - missing, quality, format issues
4. Integration - API, timeout, authentication
5. Agent Logic - wrong reasoning, out of scope

- **Process:**

- Rule-based: Pattern matching for common errors
- LLM-based: Complex error analysis with Llama 3.1

- **Output:** Error category, severity, suggested fix

#### Layer 5: Metrics Aggregation

- **Purpose:** Calculate KPIs from all monitoring data
  - **Metrics:**
    - Evaluation accuracy (overall, per-agent)
    - Drift trends over time
    - Error distribution by category
    - Error frequency and severity patterns
  - **Storage:** PostgreSQL for historical analysis
- 

## 2. Storage Layer

---

**Database:** PostgreSQL 15+ with pgvector

### Table 1: queries

**Purpose:** Central table for all queries - query\_id, query\_text, agent\_type - agent\_response, generated\_sql - status, timestamp, user\_id

## Table 2: evaluations

**Purpose:** Store evaluation results - evaluation\_id, query\_id, agent\_type - structural\_score, semantic\_score, llm\_score, final\_score - evaluation\_result (PASS/FAIL), confidence, reasoning - timestamp

## Table 3: errors

**Purpose:** Classified errors - error\_id, query\_id, evaluation\_id - error\_category, error\_subcategory - error\_message, stack\_trace, severity - frequency\_count, first\_seen, last\_seen

## Table 4: drift\_monitoring

**Purpose:** Track query drift - drift\_id, query\_id - query\_embedding (VECTOR 384) - drift\_score, drift\_classification - similarity\_to\_baseline, is\_anomaly - timestamp

## Table 5: baseline

**Purpose:** Reference for drift detection - baseline\_id, agent\_type - centroid\_embedding (VECTOR 384) - num\_queries, avg\_query\_length - common\_keywords (JSONB)

### Relationships:

```
queries (1) → (N) evaluations
queries (1) → (N) errors
queries (1) → (N) drift_monitoring
evaluations (1) → (N) errors
```

### 3. Data Flow



↓

**STEP 6: STORAGE & METRICS**

- Data stored in PostgreSQL tables
- Metrics calculated from stored data
- Results available for monitoring & analysis